

# ASSIGNMENT 2: IMPUTING METHYLATION STATUS

Barbara Engelhardt, Princeton University

out 02/23/2016; due 03/22/2016

## Background

Much of the variance in traits across individuals is regulated by genetic and epigenetic markers. While genetics is the study of DNA, which is inherited, stable through cell division, and (generally) does not change in response to environment, *epigenetics* is the study of non-genetic cellular processes that may be inherited, are stable through cell division, and may also change in response to external and internal cellular stimuli. Epigenetic markers may change within an individual over time, and may be variable across different tissues in the body. DNA methylation is possibly the most studied epigenetic modification of DNA, but we still don't have a comprehensive understanding of this epigenetic marker. In vertebrates, DNA methylation occurs by adding a methyl group to the fifth carbon of a cytosine (C) residue, mainly in the context of neighboring cytosine (C) and guanine (G) nucleotides in the genome (*CpG sites*). DNA methylation has been implicated to play an important functional role in the cell, including involvement in DNA replication and gene transcription, with substantial downstream associations with development, aging, and cancer.

The important roles of that methylation play in cellular processes imply that characterizing genome-wide DNA methylation patterns may prove instrumental to a better understanding of the regulatory mechanisms of this epigenetic phenomenon (Laird [2010]). Fortunately, recent advances in methylation-specific microarray and sequencing technologies have enabled the assay of DNA methylation patterns genome-wide at single base-pair resolution (Laird [2010]).

The current gold standard to single-site DNA methylation quantification is whole-genome bisulfite sequencing (WGBS), which quantifies DNA methylation levels at  $\sim 26$  million (out of 28 million total) CpG sites in the human genome. However, WGBS suffers from the drawback that it is prohibitively expensive (Laird [2010]). As an alternative, methylation microarrays, and the Illumina HumanMethylation 450K Beadchip in particular, measure bisulphite-treated DNA methylation levels at  $\sim 482,000$  pre-selected CpG sites, genome-wide (Zhang et al. [2013]). However, this array assays less than 2% of total CpG sites in the human genome, and the probes are biased to gene regions and CpG islands/clusters (CGIs). The purpose of this project is to develop quantitative methods to predict (*impute*) methylation status at unassayed CpG sites across the genome.

## Project definition

Your goal is to use a data set consisting of 33 reference samples with whole genome bisulfite sequence (WGBS) data (Ziller et al. [2013]) to predict (impute) the methylation levels for a sam-

ple that only has a small subset of observed methylation levels (i.e., only those CpG sites on the Illumina 450K array). We have supplied you with a reference dataset separated by chromosome for 5 chromosomes (1, 2, 6, 7 and 11). Feel free to **only use the data from chromosome 1 for this project**. Analyses on other chromosomes is optional. All of the necessary information (data, readme, this document, and other demos) is located in the github repo:

- [https://github.com/bjo/methylation\\_imputation](https://github.com/bjo/methylation_imputation)

The reference data is contained in the train.bed files. The first 3 columns of these files contain the chromosome number and position numbers corresponding to the genomic locations of each of the CpG sites, while the fourth column indicates whether the site is on the + strand or the – strand. The following 33 columns correspond to the 33 reference samples. These values,  $\beta \in [0, 1]$ , are the fraction of reads from this sample at the corresponding genomic location that are methylated. In general, most CpG sites have values of  $\beta$  that are greater than 0.7, but this will not be the case in some regions of the genome.

Your task in this project is to use the information contained in the reference (training) data file to *impute*, or predict, the  $\beta$  values for every CpG site (position) that is in the sample\_partial file that contains a value of nan. These sites will be indicated by a value of 0 in the final column of the sample\_partial file. The ‘ground truth’ for the sample\_partial file containing the  $\beta$  values for every CpG site is contained in the sample\_full file. You can compare the quality of your predictions against the known  $\beta$  values. **Do not use these ground truth values in the imputation step**, as the purpose of this exercise is to predict these values from a subset of the data and the reference sequences.

You should implement multiple methods for predicting the methylation levels at each CpG site that fits a model to the reference data set. You should carefully design the set of *features* that you will use to predict methylation values at each position. These could include, the  $\beta$  values at that position for the 33 reference samples, the values of the nearby positions in the sample\_partial.bed file, and the distance to those neighboring positions in bases. For the methods, feel free to use the generalized linear regression models we have discussed in class as well as other methods mentioned in our text books, methods described in the scientific literature, or methods implemented in software packages. You may also use more sophisticated classifiers (see *Extensions*). You may also consider using some type of regularization or sample weighting as well. Finally, you should evaluate the classifiers you apply to this problem using, at a minimum,  $r^2$  and root mean squared error (RMSE) to compare the imputed levels against the true methylation levels.

Essential to any data analysis task is the interpretation of the results. What features were the most important for prediction, and what do these features tell us about the problem? Were some reference samples more predictive of the held out sample than others? What did the distribution of prediction errors look like: were they approximately Gaussian and zero centered, or do they have a heavy tail, or something else? Simply building a machine learning approach to solve the problem does not constitute a data analysis; recovering and characterizing signal from these results does.

## Deliverables

Your deliverables for this project include:

- A four page (not including citations) summary of the project work, which should contain (as described in the Example project write up on Piazza):
  - A title, authors' names, and abstract for the project;
  - an introduction to the problem being addressed;
  - a brief description of the data;
  - a description of the methods developed and used, and how they were fit using reference data;
  - a presentation of the results of the methods applied to the test data;
  - a discussion of the results, including specific examples of single CpG sites or features that highlight the behavior of the imputation models;
  - a short summary and conclusion, including extensions that you believe would be particularly valuable based on the results;
  - a *complete* bibliography to support the methylation databases, feature selection, prediction methods, code bases, and related work that are relevant to your project.
- **please make your analysis code available to us: put a single file in the same dropbox folder with the same name as your PDF project writeup (different extension).**

Please use the  $\text{\LaTeX}$  template we have provided for you. Put your PDF write up of the project into

[https://dropbox.cs.princeton.edu/COS424\\_S2016/Assignment2](https://dropbox.cs.princeton.edu/COS424_S2016/Assignment2)

by midnight on the assignment due date, with the file name <author1PUID>\_<author2PUID>\_hw2.pdf.

Please only submit one PDF per pair of authors.

We strongly recommend *writing as you go* for this homework, which means starting to write the project report as you are downloading and analyzing the data. With that said, you should also avoid speculative writing, and only write results once you have them.

## Extensions

If you would like to extend this assignment to more interesting grounds after first completing the basic deliverables for the project, you could consider the following:

- *Extend the data set:* The data here represents a single chromosome; there are 4 additional chromosomes that you may wish to work with. Moreover, there are a number of publicly available WGBS data sets in public data bases; in particular, Gene Expression Omnibus appears to have over 60 entries corresponding to public human (*Homo sapiens*) WGBS data as of this writing. Processing and incorporating these data sets as additional

reference data—including ones you personally compile—and releasing these data with appropriate permissions would be worthwhile.

- *More interesting features*: while we have only asked you to predict methylation levels using a set of reference sequences, there are many extensions to this to consider, including features involving:
  - *Neighboring CpG sites*: there is often correlation in methylation levels among neighboring CpG sites
  - *Genetic context*: methylation levels are affected by whether the CpG site is in a CGI, intronic, or exonic regions
  - *Cis-regulatory elements*: methylation levels are correlated with certain related regulatory elements (Zhang et al. [2013])
  - *Genomic sequence properties*: methylation levels are correlated with some properties of DNA sequence, including proportions of C and G nucleotides, sequence conservation, and signatures of selection
- *More complex classifiers*: there are a number of more sophisticated statistical models that might be used for this task, including, e.g., Gaussian processes ([Roberts et al., 2013, Rasmussen, 2006]), regularized beta generalized linear models, or something of your own design that might identify latent structure in the data that is predictive of methylation levels at specific CpG sites.
- *Confidence intervals*: Imputation may be much easier for some CpG sites than for others; quantifying uncertainty in prediction values is an interesting extension that would be useful for downstream tasks.
- *Better evaluation metrics*: the downstream use of imputed methylation levels is to test for association with traits, such as cancer status or type II diabetes. What are the best ways to evaluate your predictions with this downstream task in mind? Can you improve model evaluation using cross validation instead of our training and test sets?
- *Additional types of problems*: Can you use these data to predict when two CpG positions will have different values, or test for differences in methylation status?

## Resources

There is a large literature on epigenetic markers. Many current methods for imputing methylation levels at specific CpG sites involve fairly simple classification methods and large numbers of features or reference data sets. Review some of this literature to get ideas on ways to create good imputation methods.

## References

- Peter W Laird. Principles and challenges of genomewide DNA methylation analysis. *Nature reviews. Genetics*, 11(3):191–203, March 2010. ISSN 1471-0064. doi: 10.1038/nrg2732. URL <http://www.ncbi.nlm.nih.gov/pubmed/20125086>.
- Carl Edward Rasmussen. Gaussian processes for machine learning. 2006.
- S Roberts, M Osborne, M Ebdon, S Reece, N Gibson, and S Aigrain. Gaussian processes for time-series modelling. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 371(1984):20110550, 2013.
- Weiwei Zhang, Tim D Spector, Panos Deloukas, Jordana T Bell, and Barbara E Engelhardt. Predicting genome-wide dna methylation using methylation marks, genomic position, and dna regulatory elements. *arXiv preprint arXiv:1308.2134*, 2013.
- Michael J. Ziller, Hongchang Gu, Fabian Müller, Julie Donaghey, Linus T.-Y. Tsai, Oliver Kohlbacher, Philip L. De Jager, Evan D. Rosen, David a. Bennett, Bradley E. Bernstein, Andreas Gnirke, and Alexander Meissner. Charting a dynamic DNA methylation landscape of the human genome. *Nature*, pages 1–5, August 2013. ISSN 0028-0836. doi: 10.1038/nature12433. URL <http://www.nature.com/doifinder/10.1038/nature12433>.